



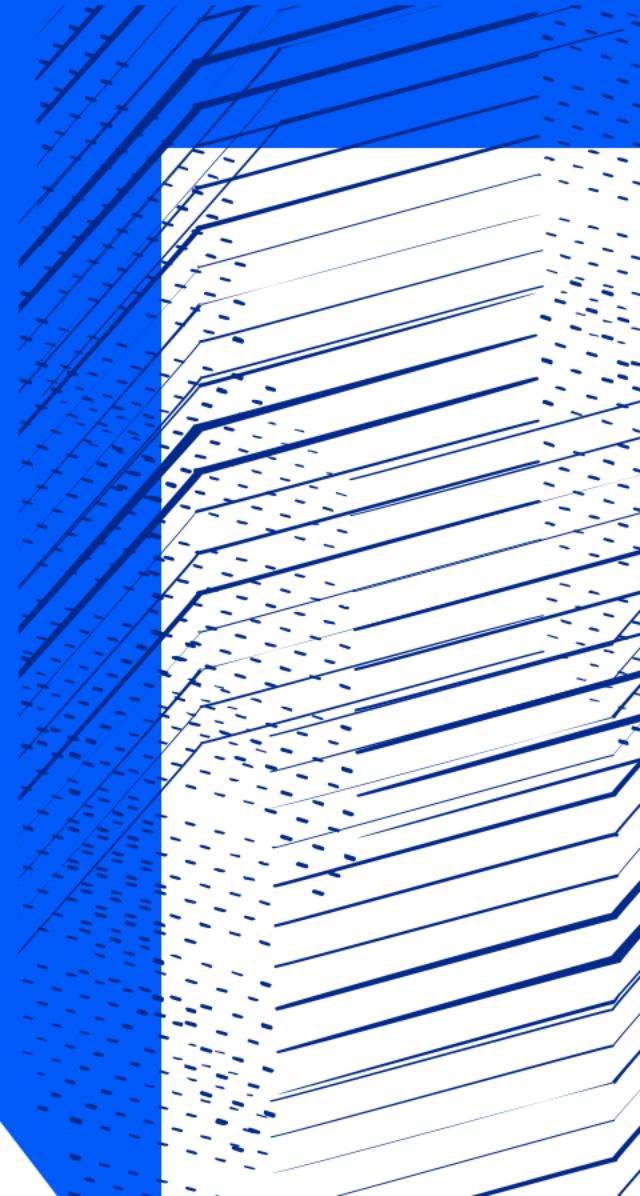
Science and
Technology
Facilities Council

Machine Learning for Science

*Generative Adversarial
Networks*

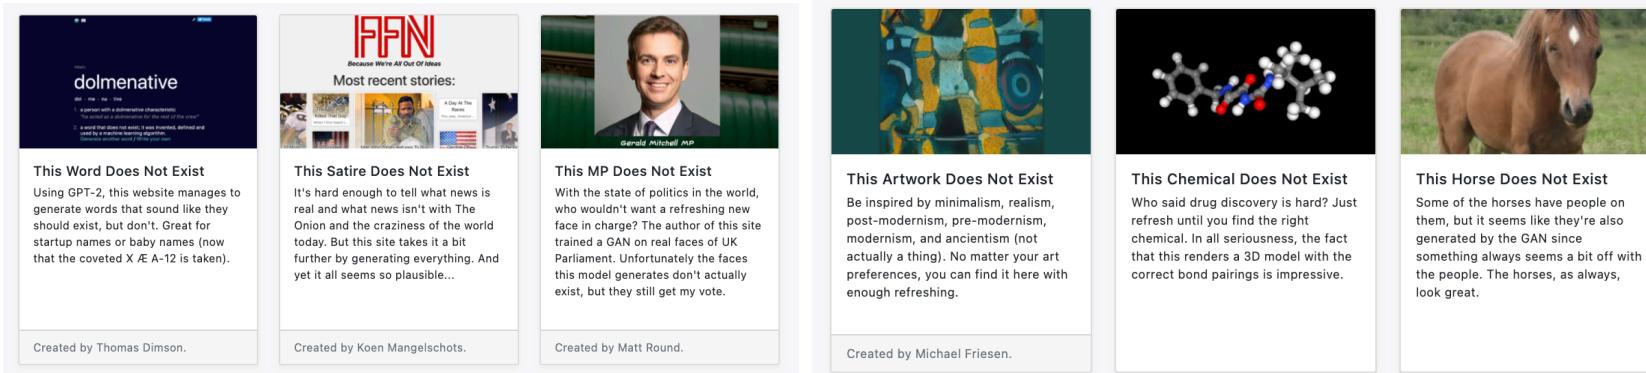
Kuangdai Leng

Scientific Machine Learning (SciML) Group
Rutherford Appleton Laboratory

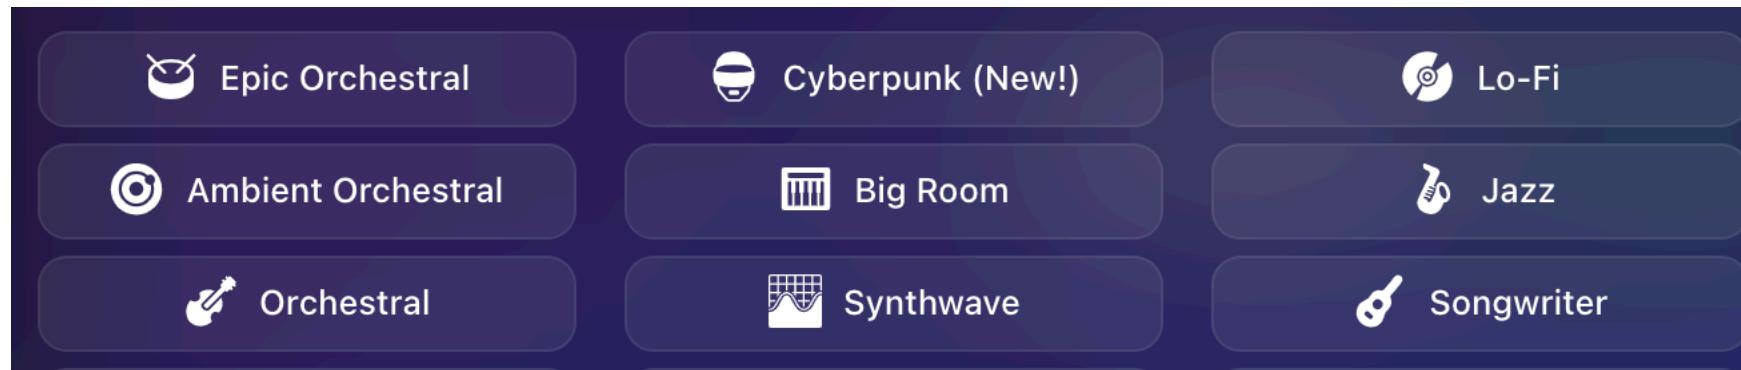


Introduction

- Generative Adversarial Networks or GANs prove to be one of the most exciting ideas in deep learning
- GANs have been used to generate *diverse realistic comprehensive* contents

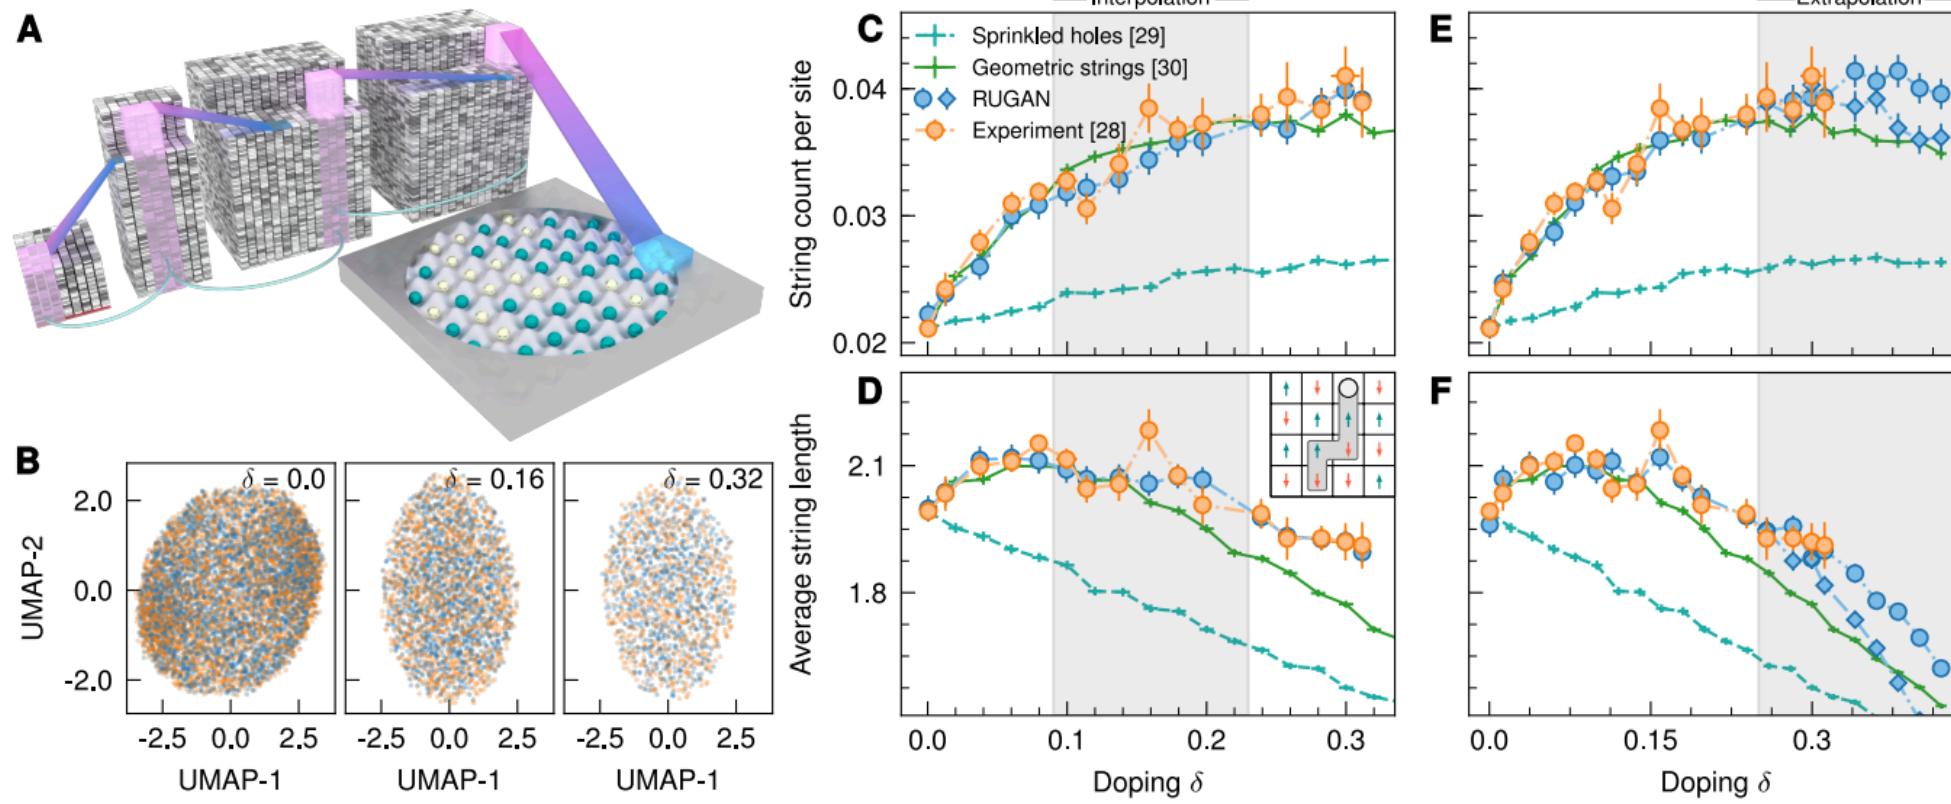


This X Does Not Exist
By Style-based GANs



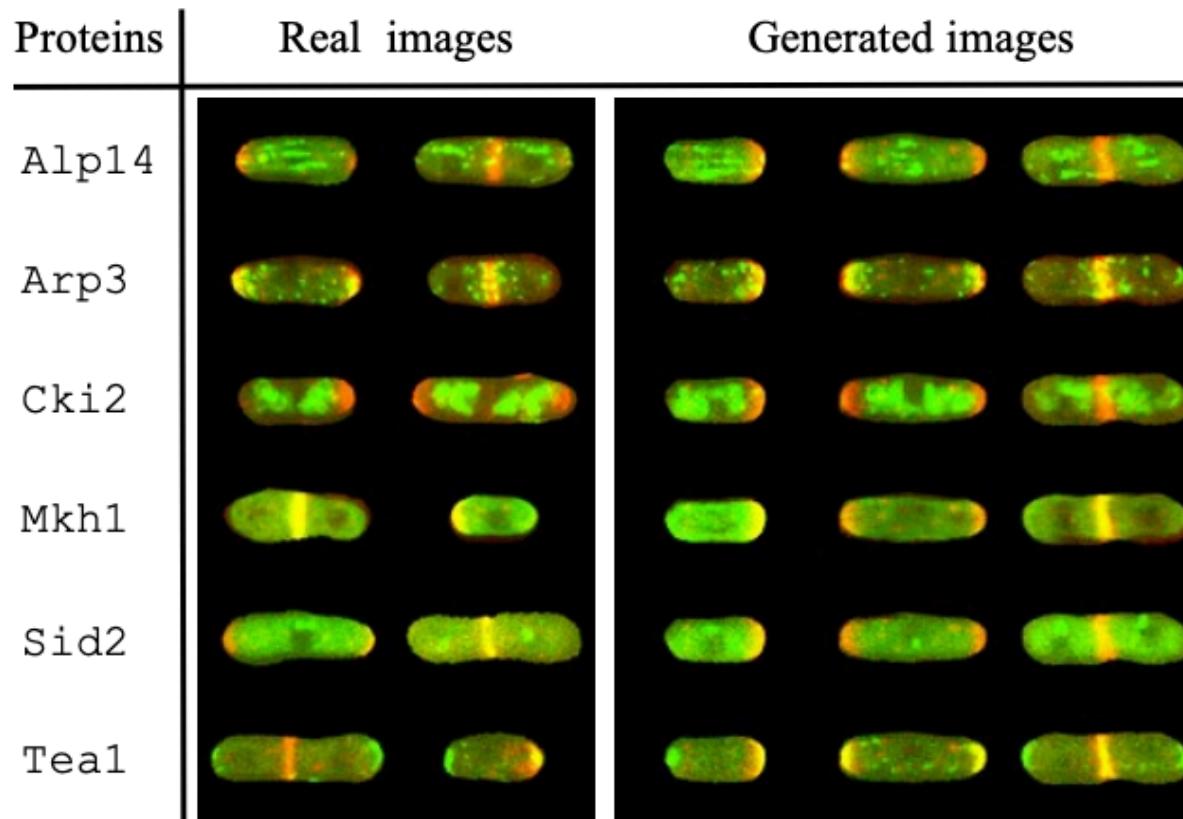
Scientific applications

- Regressive upscaling GANs to produce snapshots of a doped two-dimensional Fermi-Hubbard model (Casert et al., 2020)



Scientific applications

- GANs for Biological Image Synthesis of cells imaged by fluorescence microscopy (Osokin et al., 2017)



In this lecture

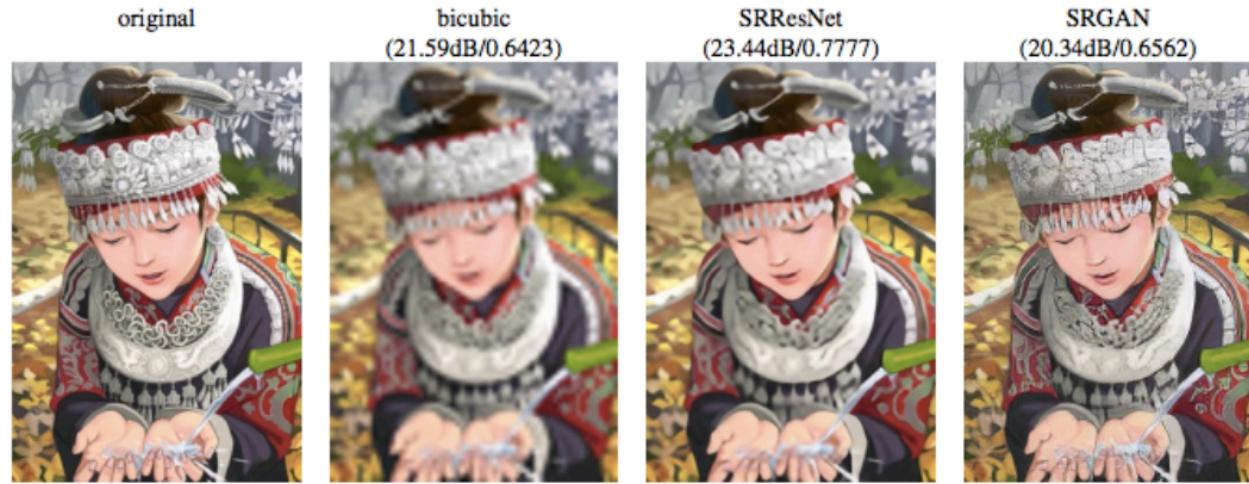
- Generative Models and Maximum Likelihood
- How GANs work and challenges
- GANs and VAEs: comparison and combination
- Summary
- Practicals: GANs for MNIST and Inelastic Neutron Scattering

Following the structure of

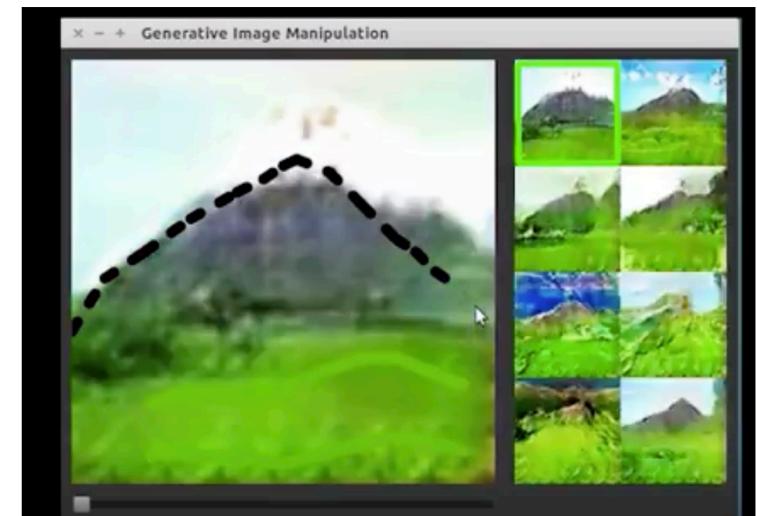
NIPS 2016 Tutorial: Generative Adversarial Networks, Ian Goodfellow, 2016

Why generative models?

- High-dimensional probability distributions
- Realistic content generation (fake is good or even better, e.g., art)
- Data expansion and missing labels (semi-supervised)
- Reinforcement learning
- Inspiring human beings



SRGAN
Ledig et al. 2016



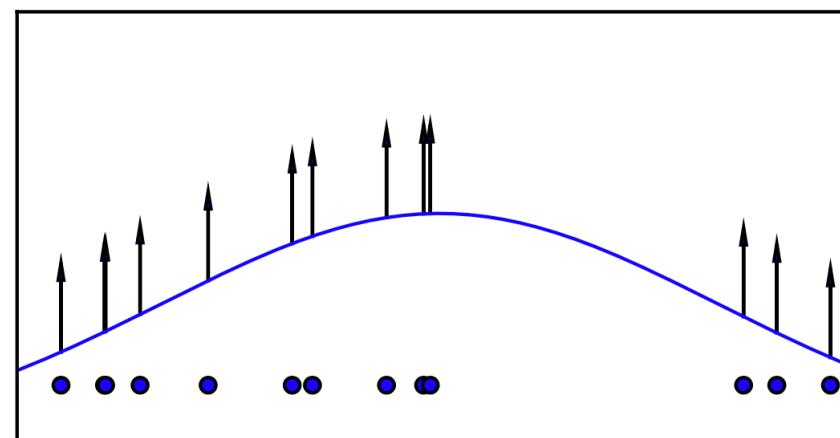
Interactive GAN (iGAN)
Zhu et al. 2016

Maximum Likelihood Estimate

- Maximum Likelihood
- Minimising KL-divergence between data and model distributions

$$\begin{aligned}\boldsymbol{\theta}^* &= \arg \max_{\boldsymbol{\theta}} \prod_{i=1}^m p_{\text{model}}(\mathbf{x}^{(i)}; \boldsymbol{\theta}) \\ &= \arg \max_{\boldsymbol{\theta}} \log \prod_{i=1}^m p_{\text{model}}(\mathbf{x}^{(i)}; \boldsymbol{\theta}) \\ &= \arg \max_{\boldsymbol{\theta}} \sum_{i=1}^m \log p_{\text{model}}(\mathbf{x}^{(i)}; \boldsymbol{\theta}).\end{aligned}$$

$$\boldsymbol{\theta}^* = \arg \min_{\boldsymbol{\theta}} D_{\text{KL}}(p_{\text{data}}(\mathbf{x}) \| p_{\text{model}}(\mathbf{x}; \boldsymbol{\theta})).$$

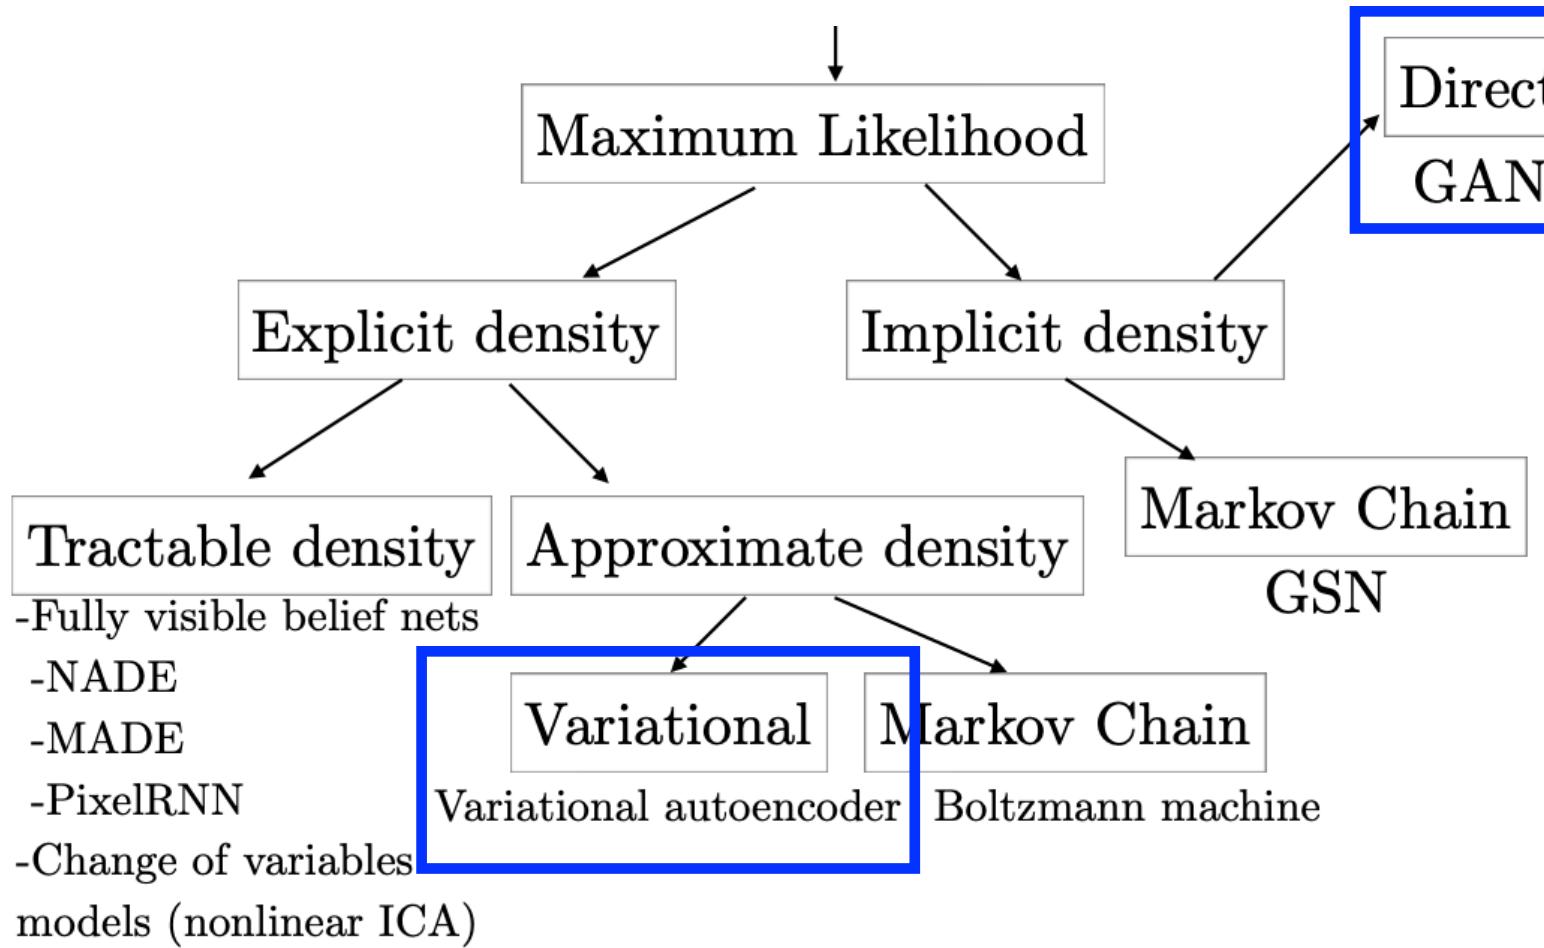


$p_{\text{data}}(\mathbf{x})$ is unknown

$$\boldsymbol{\theta}^* = \arg \max_{\boldsymbol{\theta}} \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log p_{\text{model}}(\mathbf{x} | \boldsymbol{\theta})$$

Maximum Likelihood Estimate

- Taxonomy of generative models based on Maximum likelihood

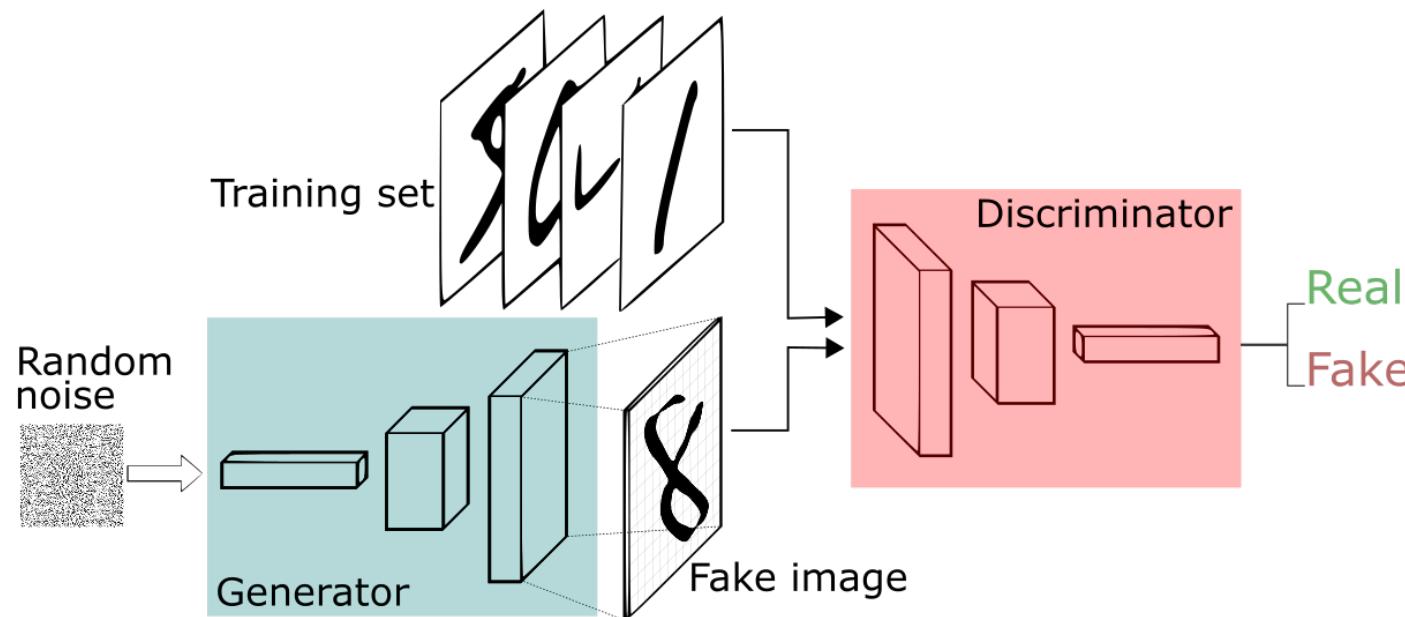


In this lecture

- Generative Models and Maximum Likelihood
- How GANs work and challenges
- GANs and VAEs: comparison and combination
- Summary
- Practicals: GANs for MNIST and Inelastic Neutron Scattering

How GANs work

- A GAN contains a generator and a discriminator competing in a two-player game
 - **Generator:** generates fake images from a latent sampling, $G(z) \rightarrow x'$, and tries to fool the discriminator that these images are real
 - **Discriminator:** tries to distinguish between fake and real images, $D(x) \rightarrow 0/1$



They can be regard as collaborators because they share information completely: the discriminator is more like a teacher than an adversary (Goodfellow, 2016)

How GANs work

- Discriminator Loss

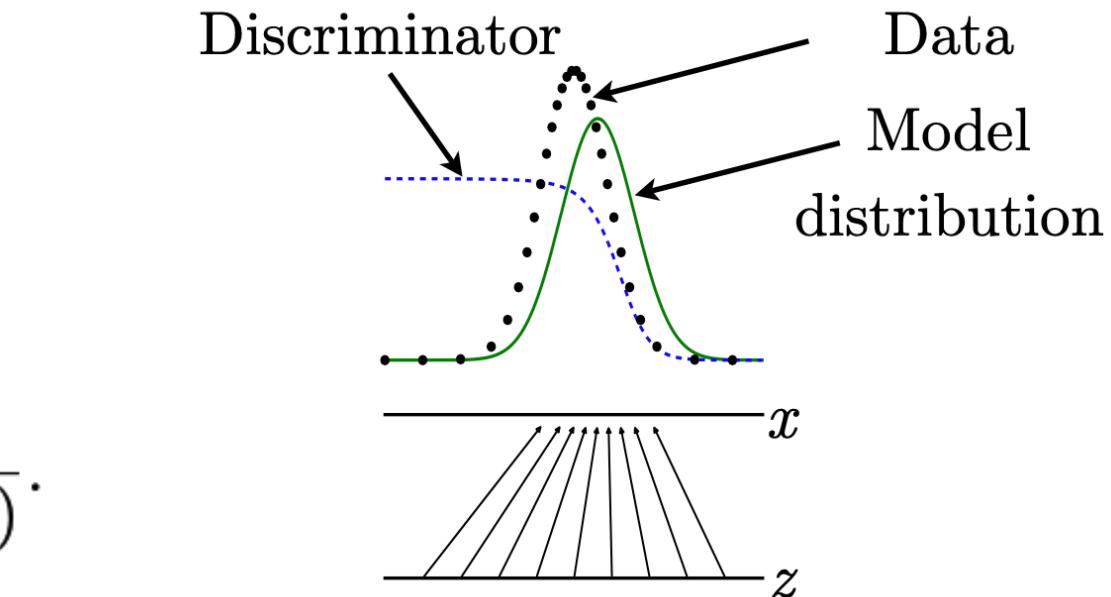
$$J^{(D)}(\boldsymbol{\theta}^{(D)}, \boldsymbol{\theta}^{(G)}) = -\frac{1}{2} \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log D(\mathbf{x}) - \frac{1}{2} \mathbb{E}_{\mathbf{z}} \log (1 - D(G(z))).$$

Disc decision
for real images Disc decision
for fake images

- The ratio of densities

$$\frac{\delta}{\delta D(\mathbf{x})} J^{(D)} = 0.$$

→ $D^*(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_{\text{model}}(\mathbf{x})}.$



How GANs work

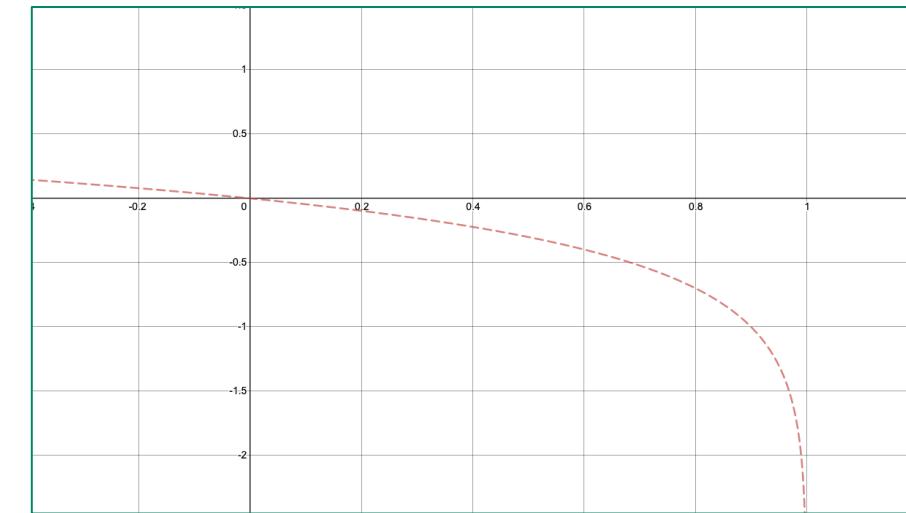
- Generator Loss: zero-sum game (simplest case)

$$J^{(G)} = -J^{(D)}.$$

- The minimax of a value function

$$\boldsymbol{\theta}^{(G)*} = \arg \min_{\boldsymbol{\theta}^{(G)}} \max_{\boldsymbol{\theta}^{(D)}} V \left(\boldsymbol{\theta}^{(D)}, \boldsymbol{\theta}^{(G)} \right).$$

where $V \left(\boldsymbol{\theta}^{(D)}, \boldsymbol{\theta}^{(G)} \right) = -J^{(D)} \left(\boldsymbol{\theta}^{(D)}, \boldsymbol{\theta}^{(G)} \right).$



- Minimization in an outer loop
- Maximization in an inner loop
- Useful for theoretical analysis
- Impractical because of vanishing gradient of generator
($\log(1 - x)$ is flat near the fake end)

How GANs work

- Generator Loss: heuristic game

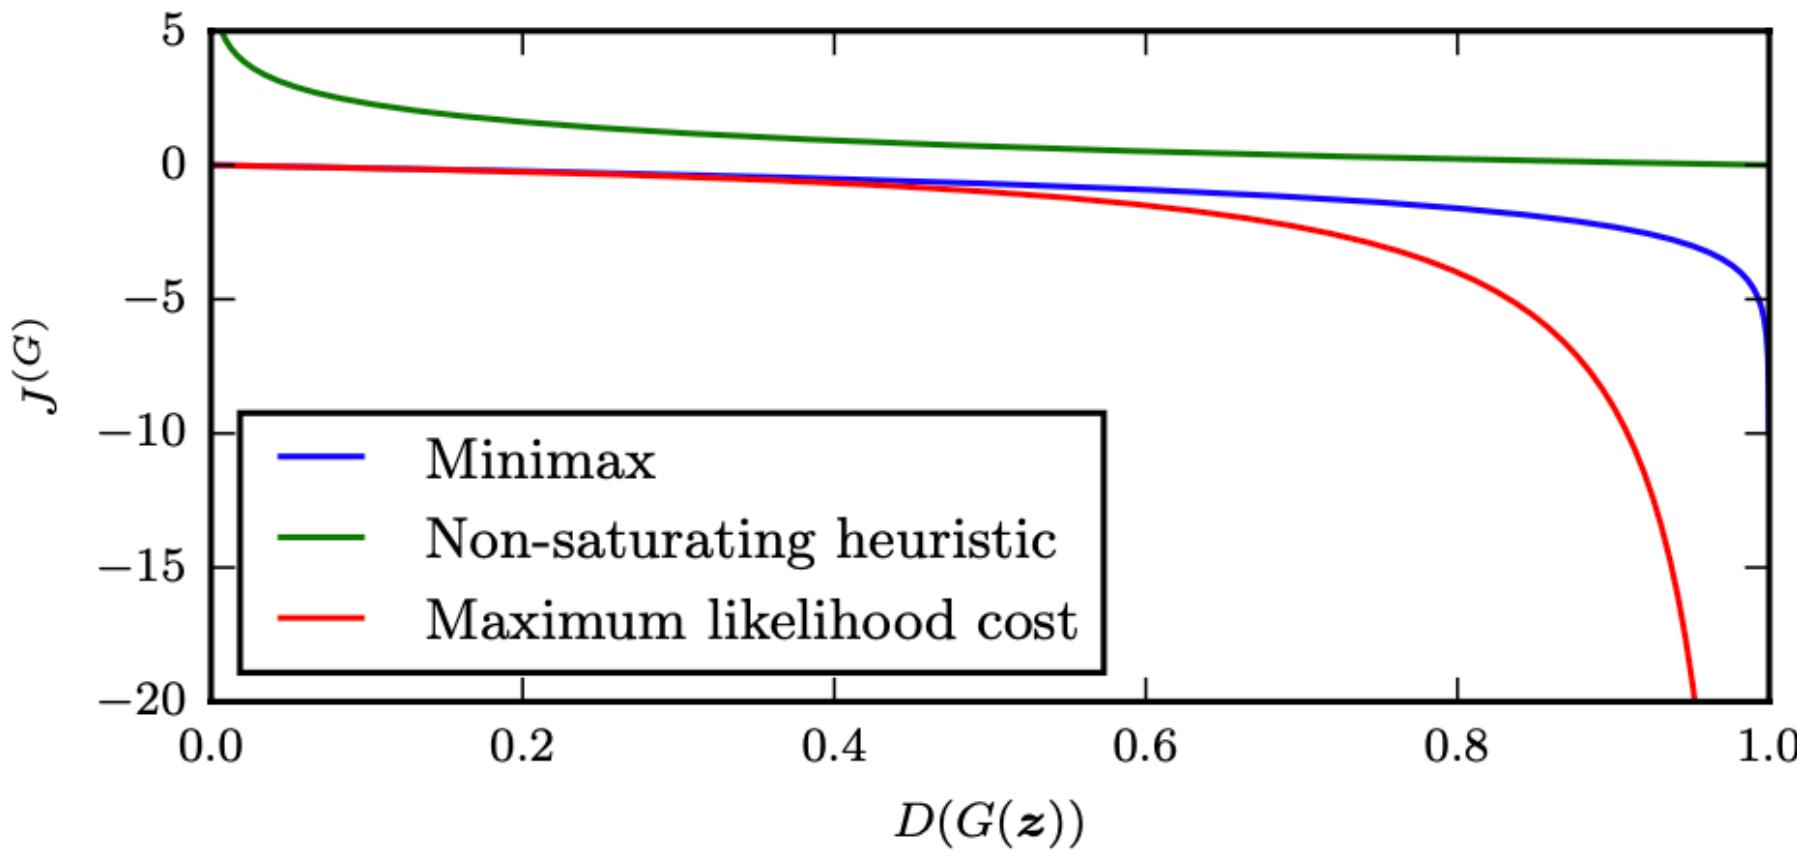
$$J^{(G)} = -\frac{1}{2} \mathbb{E}_z \log D(G(z))$$

- In the minimax game, the generator minimizes the log-probability of the discriminator being correct
- In this game, the generator maximizes the log-probability of the discriminator being mistaken
- Generator Loss: maximum likelihood

$$J^{(G)} = -\frac{1}{2} \mathbb{E}_z \exp(\sigma^{-1}(D(G(z))))$$

How GANs work

- Generator Loss comparisons



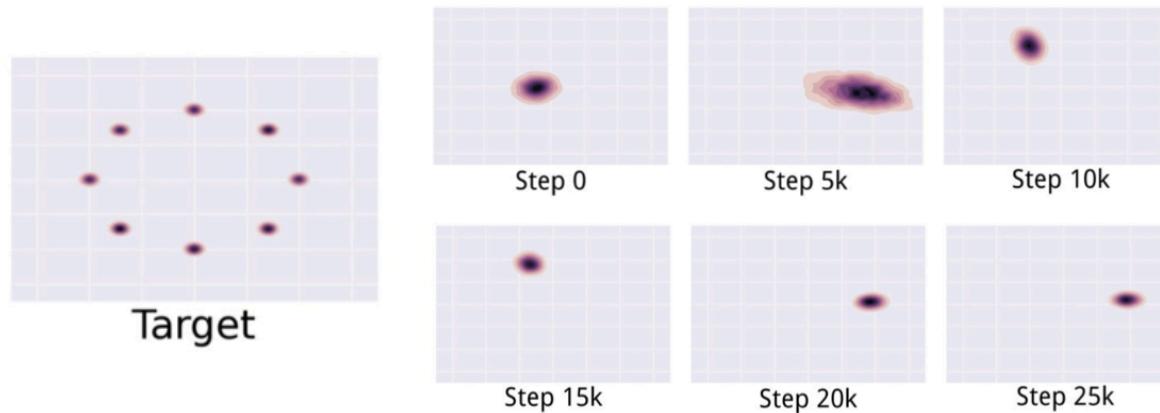
- Variance distribution
- Variance reduction

Challenges: Nash Equilibrium

- G and D are competing; if one stops evolving, both will stop (vanishing gradient)
 - What if G is much stronger than D ?
 - What if D is much stronger than G ?
- Balancing G and D is not straightforward because it depends on $p_{\text{data}}(x)$
- What we can do:
 - Loss parameterisation
 - Model architecture
 - Asynchronous training (one exercise in the notebook)

Challenges: Mode Collapse

- Mode collapse is the most significant challenge for GANs
- The generator decodes every z into one or one type of x

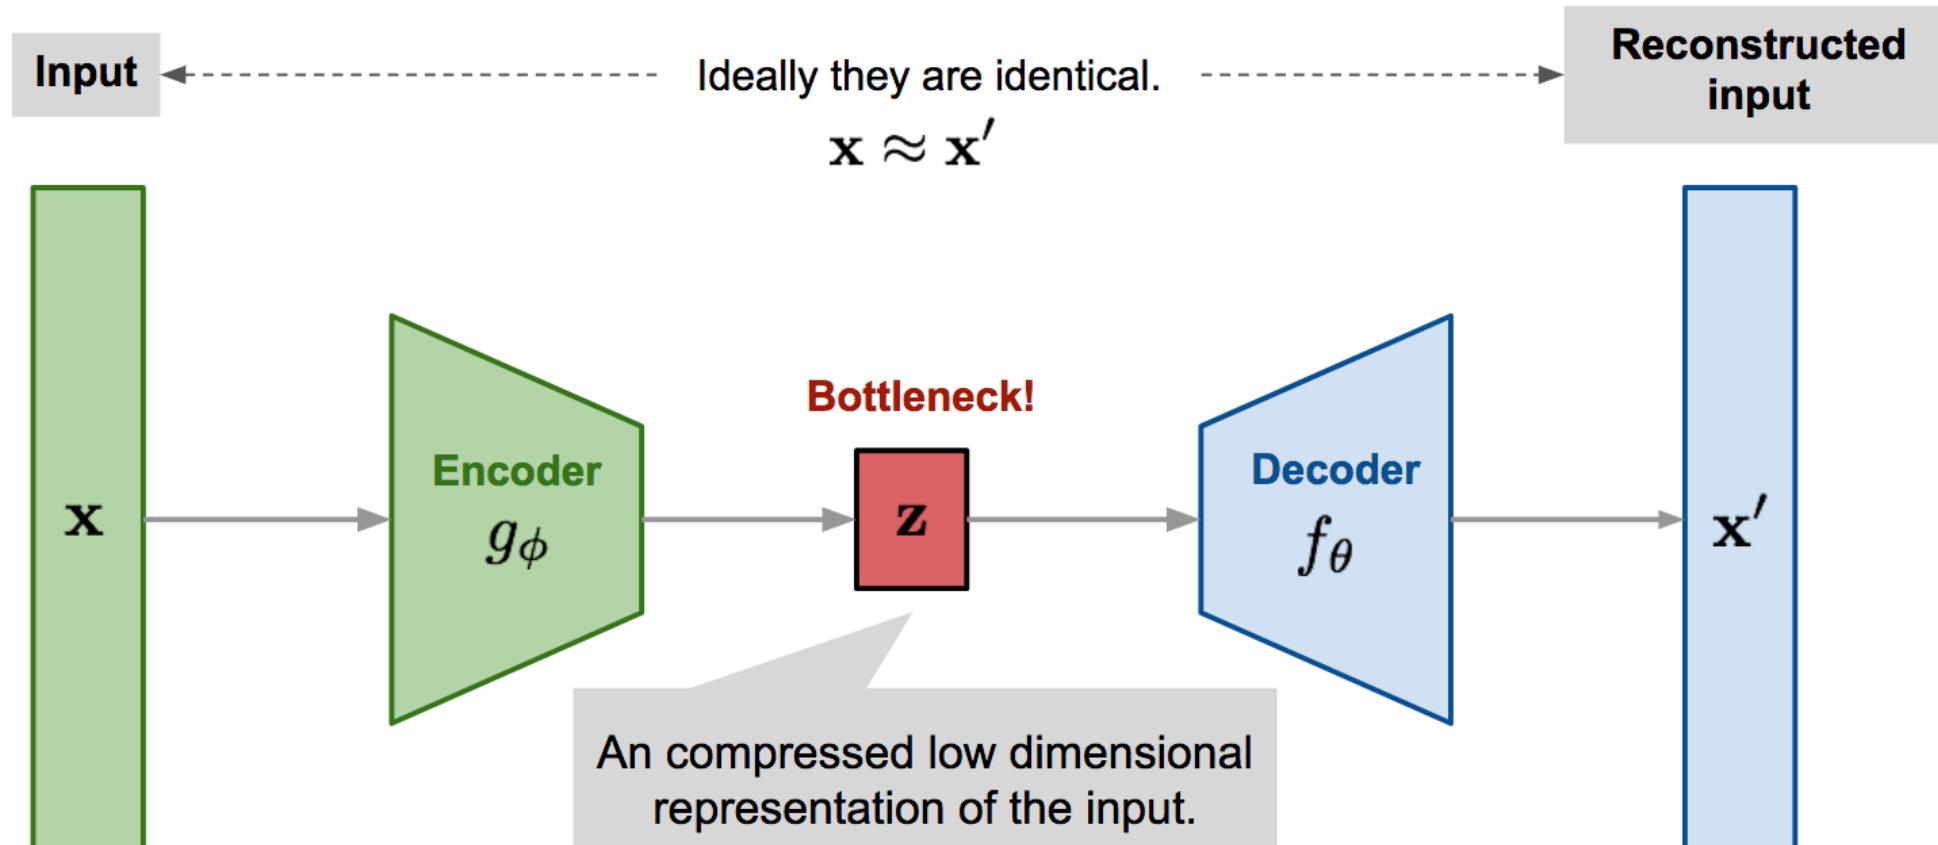


- This is an unsolved problem yet with some solutions
 - Mini-batch features (Salimans et al., 2016)
 - Unrolled GANs (Metz et al., 2016)
 - Wasserstein loss (Arjovsky et al., 2017)
 - Discriminator loss: $D(x) - D(G(z))$; Generator loss: $D(G(z))$

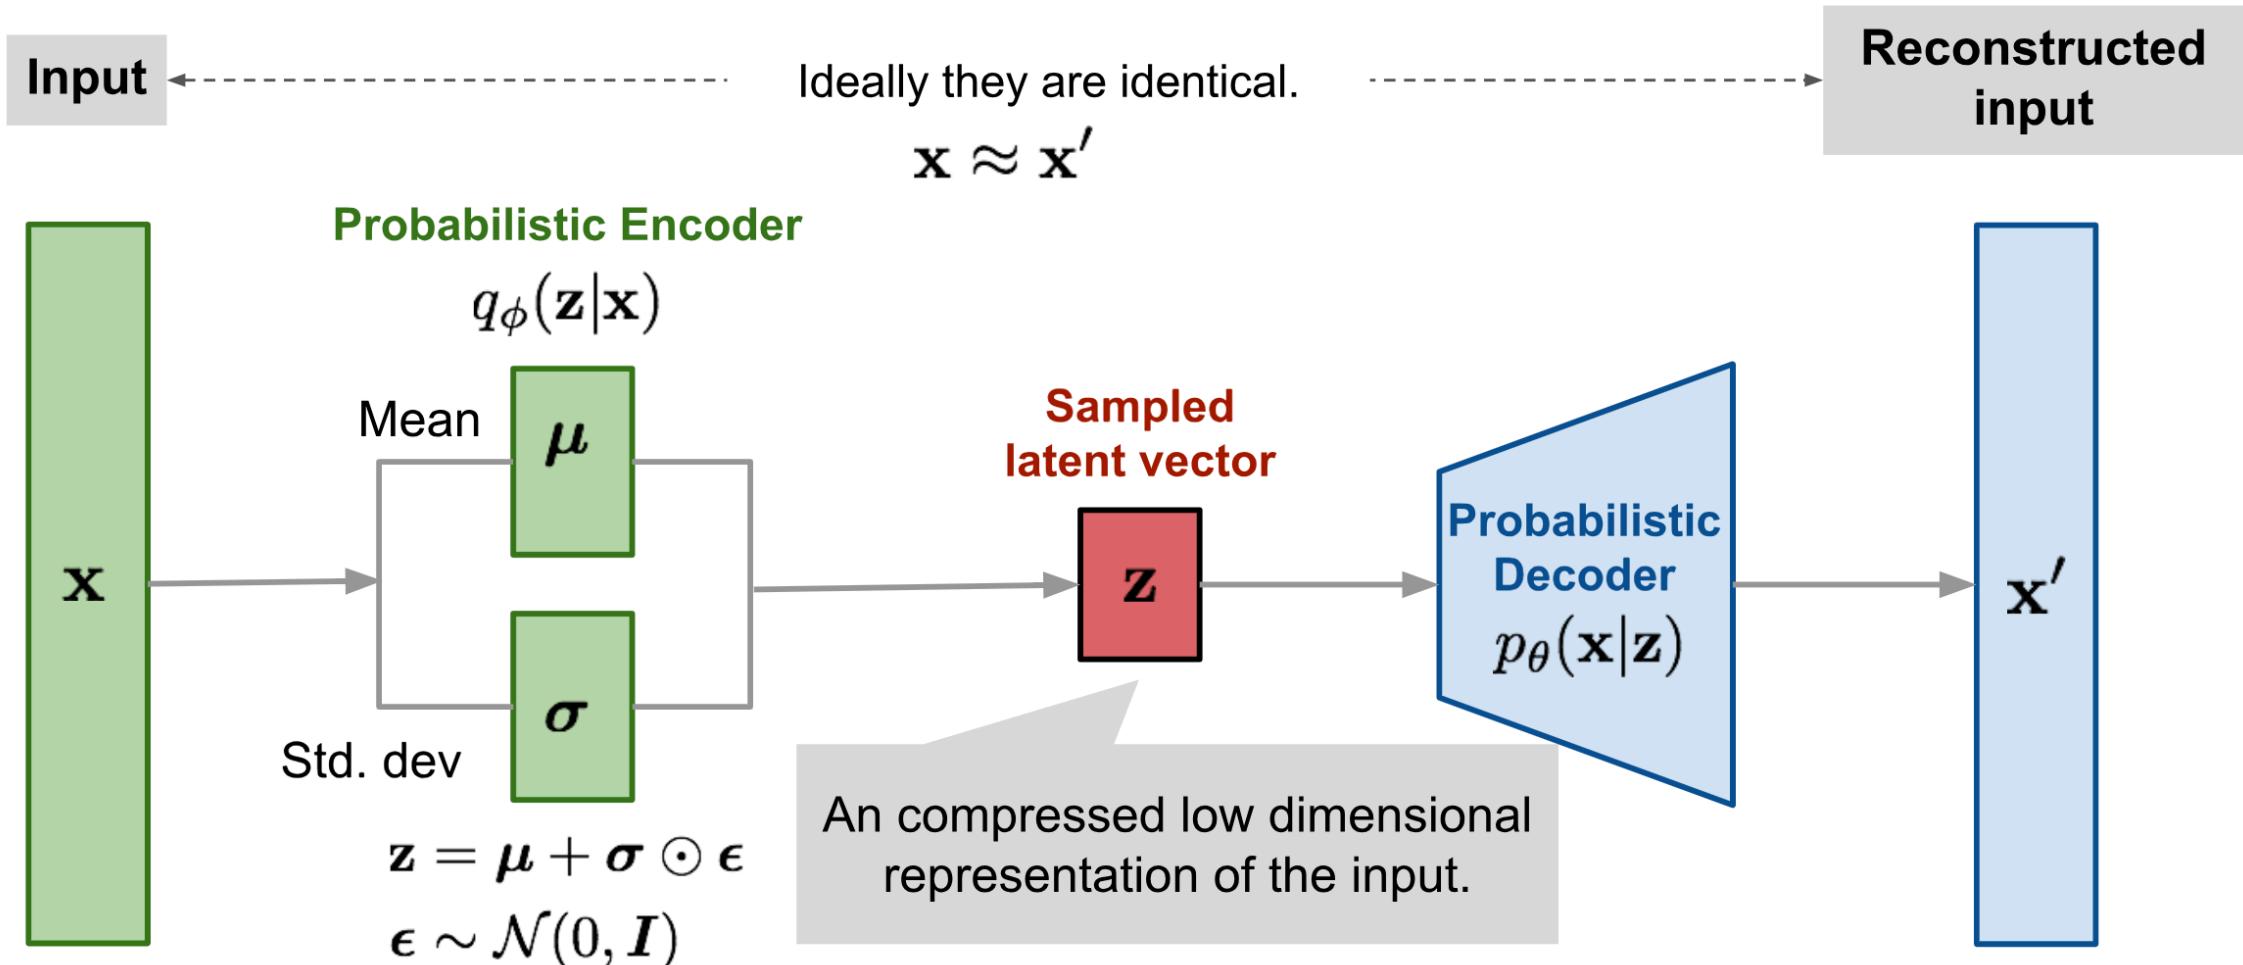
In this lecture

- Generative Models and Maximum Likelihood
- How GANs work and challenges
- GANs and VAEs: comparison and combination
- Summary
- Practicals: GANs for MNIST and Inelastic Neutron Scattering

Autoencoder (AE) and Variational Autoencoder (VAE)



Autoencoder (AE) and Variational Autoencoder (VAE)



Variational inference

- True generative process (e.g., PDE for a physics problem)

$$z \quad \Rightarrow \quad p(x|z) \quad \Rightarrow \quad x$$

- Given x (image data), we want to infer $p(z|x)$, the *posterior probability*

$$p(z|x) = \frac{p(x|z)p(z)}{p(x)} = \frac{p(x|z)p(z)}{\int p(x|z')p(z') dz'} \quad \text{intractable (MC)}$$

- Variational inference: use some distribution $q(z|x)$ to approximate $p(z|x)$

$$q(z|x) \sim p(z|x) \quad \Leftarrow \quad \min \text{KL}(q(z|x), p(z|x))$$

Variational inference

- Kullback–Leibler divergence: “distance” between two distributions

$$\text{KL}(q, p) = - \mathbb{E}_q \log \frac{p}{q} \quad \Rightarrow \quad \text{KL}(q, p) \geq 0 \quad \text{KL}(q, p) \neq \text{KL}(p, q)$$

Don't have to use KL divergence (β -VAE)

- Variational identity

$$\text{KL}(q(z|x), p(z|x)) = \log p(x) - (-\text{KL}(q(z|x), p(z)) + \mathbb{E}_{q(z|x)} \log p(x|z))$$

- Evidence lower bound

$$\log p(x) \geq -\text{KL}(q(z|x), p(z)) + \mathbb{E}_{q(z|x)} \log p(x|z)$$

ELBO

Variational Autoencoder

- Evidence lower bound

$$\log p(x) \geq -\text{KL}(q(z|x), p(z)) + \mathbb{E}_{q(z|x)} \log p(x|z)$$



$$\begin{aligned} \log p_\theta(x) &\geq -\text{KL}(q_\phi(z|x), p(z)) \\ &\quad + \mathbb{E}_{q_\phi(z|x)} \log p_\theta(x|z). \end{aligned}$$

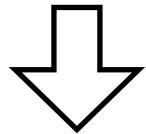
- θ : Generative model or the decoder (θ looks like g)
- ϕ : Inference model or the encoder (ϕ has an I in the middle)

Variational Autoencoder

- Maximum likelihood training

$$\underline{\mathbb{E}_{p_{\mathcal{D}}(x)} \log p_{\theta}(x)},$$

Data



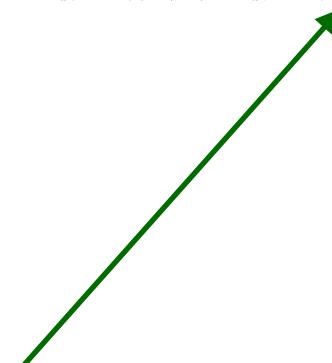
$$\max_{\theta} \max_{\phi} \mathbb{E}_{p_{\mathcal{D}}(x)} \left[\underline{-\text{KL}(q_{\phi}(z | x), p(z))} + \underline{\mathbb{E}_{q_{\phi}(z|x)} \log p_{\theta}(x | z)} \right].$$

Reconstruction error

$$q_{\phi}(z | x) = \mathcal{N}(\mu, \sigma^2)$$
$$p(z) = \mathcal{N}(0, 1)$$

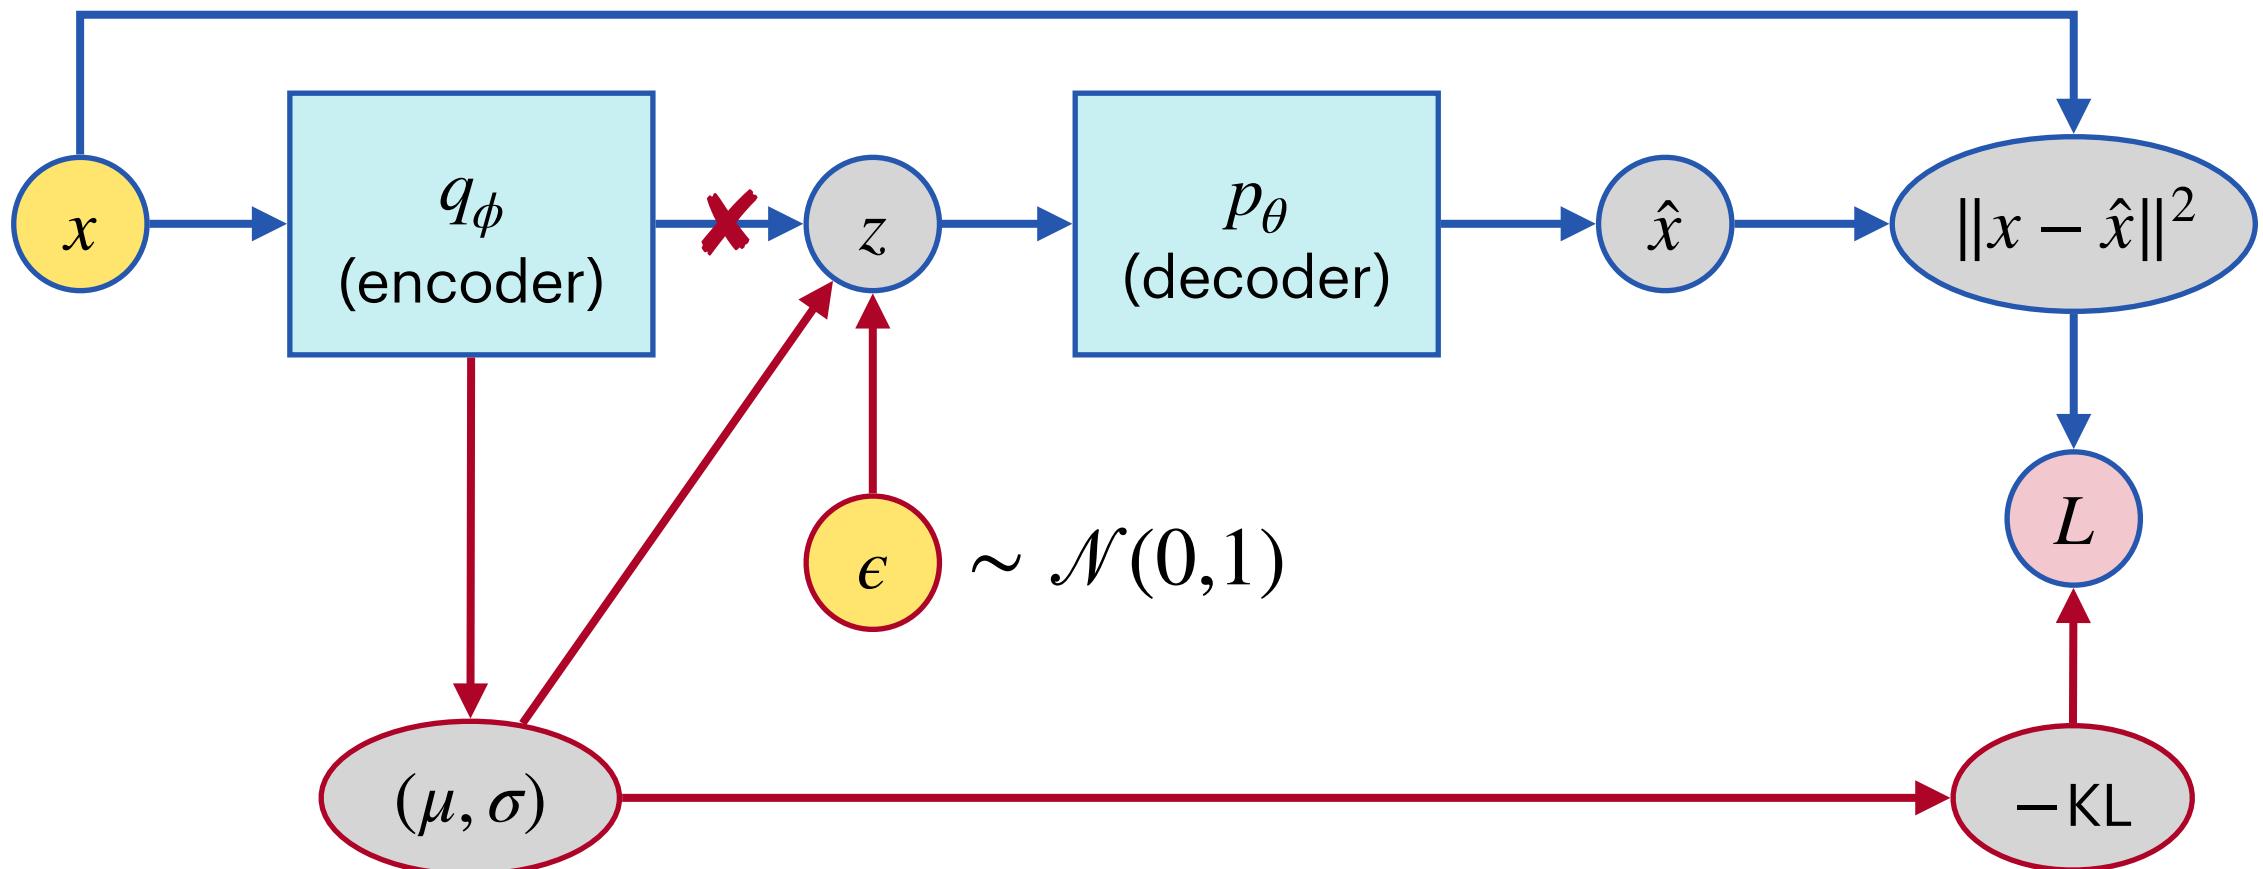
Prior

**Approx.
posterior**

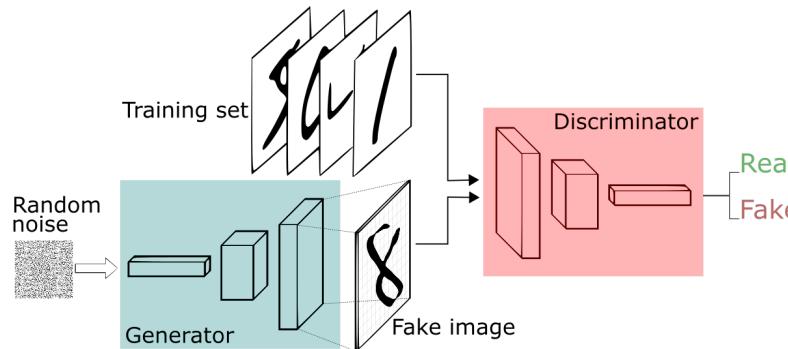


Variational Autoencoder

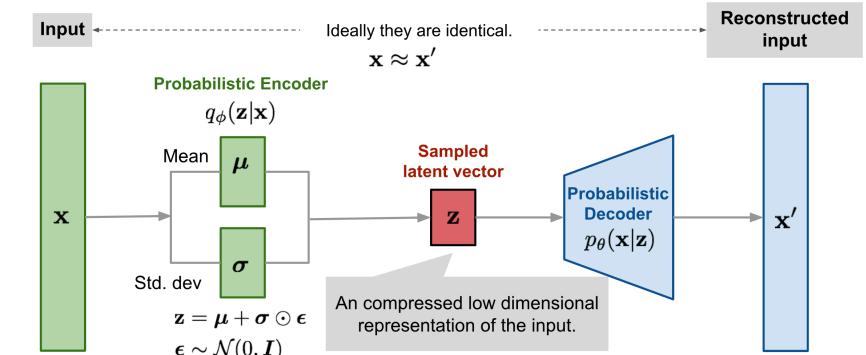
- Architecture



GANs and VAEs

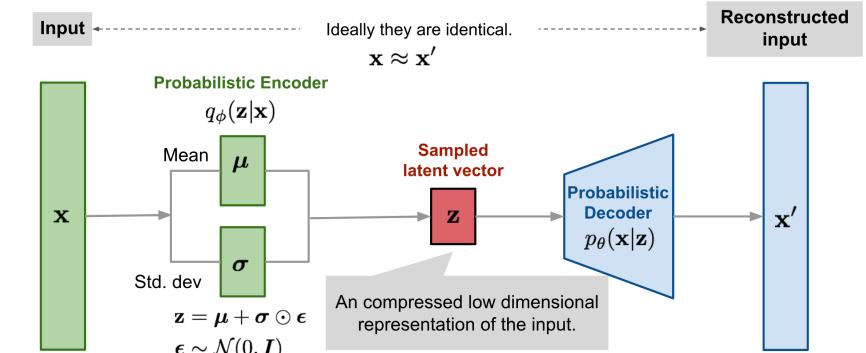
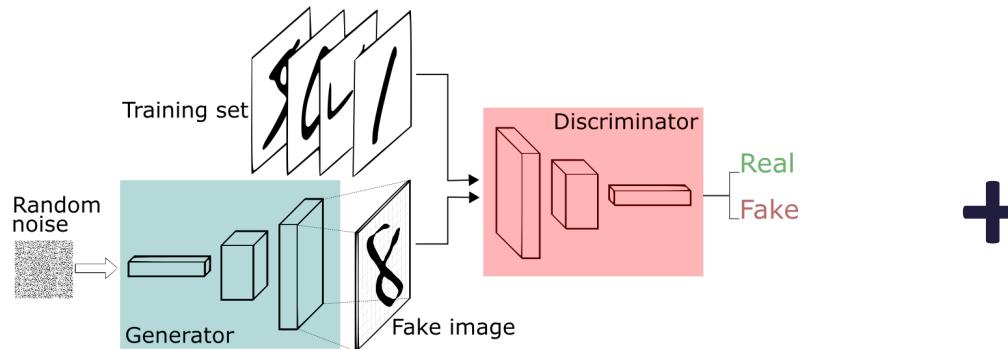


VS



	Sharp	Content generation	Vague
icGAN, infoGAN	No	Inference model	Yes
	Hard	Feature control and PCA	Easy
	Hard	Designing & training	Easy
	Implicit Indirect	Maximum likelihood Generator sees data	Explicit Direct

GANs and VAEs



- Notable efforts to combine GANs and VAEs
 - **Simple merge** (Larsen et al., 2015): VAE's decoder as GAN's generator
 - **Adversarial AE** (Makhzani et al., 2015): VAE's latent (+ seed) as input of GAN's generator
 - **Bidirectional GANs** (Donahue et al., 2016): Both VAE's encoder and decoder provide (feature, data) pairs as the input of GAN
 - **Adversarial Variational Bayes** (Lars et al., 2017): Use a GAN to compute VAE's variational loss without assuming a posterior form (most theoretically elegant)

In this lecture

- Generative Models and Maximum Likelihood
- How GANs work and challenges
- GANs and VAEs: comparison and combination
- [Summary](#)
- Practicals: GANs for MNIST and Inelastic Neutron Scattering

Summary

- GANs are powerful to generate realistic contents of nearly all types
- The generator and discriminator compete in a two-player game
- The choice of loss functions is a key issue
- GANs are difficult to train (instability)
 - Vanishing gradient
 - Mode collapse
- Lack of an inference model (important to science) is a major drawback
- Combing VAEs and GANs is a promising research direction

In this lecture

- Generative Models and Maximum Likelihood
- How GANs work and challenges
- GANs and VAEs: comparison and combination
- Summary
- **Practicals: GANs for MNIST and Inelastic Neutron Scattering**



Science and
Technology
Facilities Council

Thank you

A large, stylized graphic element in the center-right of the slide. It consists of a diagonal band of orange on the left transitioning into a blue area on the right. Within this blue area, numerous thin, light blue lines radiate outwards from the bottom-left corner, creating a dynamic, fan-like or signal-like pattern.

Science and Technology Facilities Council



@STFC_matters



Science and Technology Facilities Council